


— 4w    —○— GEMM kernel speedup    —■— GEMM share of stock time  
 — 8da4w    —●— end-to-end speedup    ·◇· attention speedup / share (GPUs with SARC attention)



model size (Llama 3.2 1B / 3B, Llama 3.1 8B), 2048-token prefill

† pre-release: unverified rows, different .pte export; drivers AMDVLK (RX 7900 XTX), RADV Mesa 26.2.3 (RX 7600)