


4w -○- GEMM kernel speedup ■ GEMM share of stock time
 8da4w ● end-to-end speedup ◇ attention speedup / share
 (GPUs with SARC attention)



model size (Llama 3.2 1B / 3B, Llama 3.1 8B), 2048-token prefill

† pre-release: unverified rows, different .pte export, AMDVLK