


Measured end-to-end: re-tuned kernels speed up real prefill by up to 1.8x

■ previous best WMMA (= 1x) ■ measured gain (colour = GPU) — paired 95 % CI ○ projection



2048-token real-text prompt, 5 interleaved repeats per build. † previous kernels were opt-in (default ran tiled).