



Prefill throughput before → after the re-tuning (measured, 2048-token prompt)

○ previous best WMMA ● re-tuned (colour = GPU)


Llama 3.2 1B



Llama 3.2 3B



Llama 3.1 8B



Prefill throughput (tokens/s, log scale)

Median of 5 interleaved repeats per build, real-text prompt. 4070 Ti 4w is flat or lower (accuracy fix); 780M 8da4w is unchanged. Previous kernels on B580, B70, 4070 Ti and Orin were opt-in.