



■ Llama 3.2 1B ■ Llama 3.2 3B ■ Llama 3.1 8B



2048-token prompt, ExecuTorch Vulkan (llama_main). Bars: ratio of medians of n=5 runs; error bars: approx. 95 % paired bootstrap CI; dashed line: 1.0x (stock).

† pre-release: unverified kernel rows (RX 7900 XTX, RX 7600);

drivers: AMDVLK (RX 7900 XTX), RADV Mesa 26.2.3 (RX 7600).

‡ Xclipse (M51): internal device, relative speedups only; pre-release rows.

4w on Xclipse (M51): kernel fix merged, end-to-end re-measurement pending; not shown.