



Llama 3.2 1B

Llama 3.2 3B

Llama 3.1 8B



2048-token prompt, ExecuTorch Vulkan (llama_main). Bars: ratio of medians of n=5 runs; error bars: approx. 95 % paired bootstrap CI; dashed line: 1.0x (stock).

† RX 7900 XTX: pre-release. SARC rows unverified (ET_VK_SARC_UNVERIFIED=1); different .pte export, so absolute tok/s are not strictly comparable; AMDVLK driver.

‡ Xclipse (M51): internal device, relative speedups only; pre-release rows.

4w on Xclipse (M51): not shown (kernel in development).