


Llama 3.1 8B · 8-bit activations + 4-bit weights



† pre-release kernel rows and a different model export (RX 7900 XTX, RX 7600)
2048-token prompt · ExecuTorch Vulkan · median of 5 runs