


Llama 3.1 8B · 8-bit activations + 4-bit weights



2048-token prompt · ExecuTorch Vulkan · median of 5 runs