


Time to first token · Llama 3.1 8B · 4-bit weights



Each row scaled to its own ExecuTorch 1.5 time

† pre-release kernel rows (RX 7900 XTX, RX 7600)
2048-token prompt · ExecuTorch Vulkan · median of 5 runs