


Time to first token · Llama 3.1 8B · 4-bit weights



† RX 7900 XTX: pre-release kernel rows, different model export
2048-token prompt · ExecuTorch Vulkan · median of 5 runs